Report to Deep Learning Applications

Lecheng WANG, Guanyu CHEN
M2A, Sorbonne University

Contents

2-a: Transfer Learning

★Question 1

A	B	C	D	E	F	G
	Function	Input	Output	um parameter	porportion	
1	Conv2d	3	64	1,792	14,714,688	sum(convolution)
2	Conv2d	64	64	36,928	123,642,856	sum(linear)
3	MaxPool2d	-	-		0.119009609	conv/linear
4	Conv2d	64	128	73,856	138,357,544	total
5	Conv2d	128	128	147,584		
6	MaxPool2d	-	-			
7	Conv2d	128	256	295,168		
8	Conv2d	256	256	590,080		
9	Conv2d	256	256	590,080		
10	MaxPool2d	-	-			
11	Conv2d	256	512	1,180,160		
12	Conv2d	512	512	2,359,808		
13	Conv2d	512	512	2,359,808		
14	MaxPool2d	-	-			
15	Conv2d	512	512	2,359,808		
16	Conv2d	512	512	2,359,808		
17	Conv2d	512	512	2,359,808		
18	MaxPool2d	-	-			
19	Linear	25,088	4,096	102,764,544		
20	Linear	4,096	4,096	16,781,312		
21	Linear	4,096	1,000	4,097,000		

Figure 1: A table to count the parameters in each layer.

This table shows nearly 90% of parameters are produced by fully connected layers.

★Question 2

Output size of last layer of VGG16 is (1,1000), each element in the vector represents the probability (or score) of the input image belonging to one of the 1000 classes.

★Question 3

Role of the ImageNet normalization:

1. **Stabilize the training procedure:** Stabilize the training procedure. Normalization reduces the chance of vanishing or exploding gradients, thus, leads to a faster and more stable convergence.
2. **Improve generalization ability:** Improve generalization ability. Normalization helps to center the data and scale it, reducing the effects of, for example, varying light conditions. This allows the model to focus on more important features like shapes and textures thus, leads to a better performance.

Setting the model to evaluation mode ensures that it remains unchanged, preventing accidental modifications. Additionally, since gradient calculations are not required, computation becomes faster. In evaluation mode, networks with dropout layers will also behave differently compared to training mode.

Question 4

We can see the different output channels have focus on different features. We use (i,j) to refer the picture in i th row and j th column. Picture (4,2) and (4,4) extract features from large, distinctly varied color blocks. Such feature extraction helps the network differentiate between the background and the main subject of the image. Picture (4,3) can be regarded as the extraction of textures in the finer details. Picture (2,4) high lights the brighter areas within the image.

It is difficult for the human brain to fully understand the meaning of each channel, but there is no doubt that each of these channels extracts some type of feature of the image, and these features are useful for later operations.

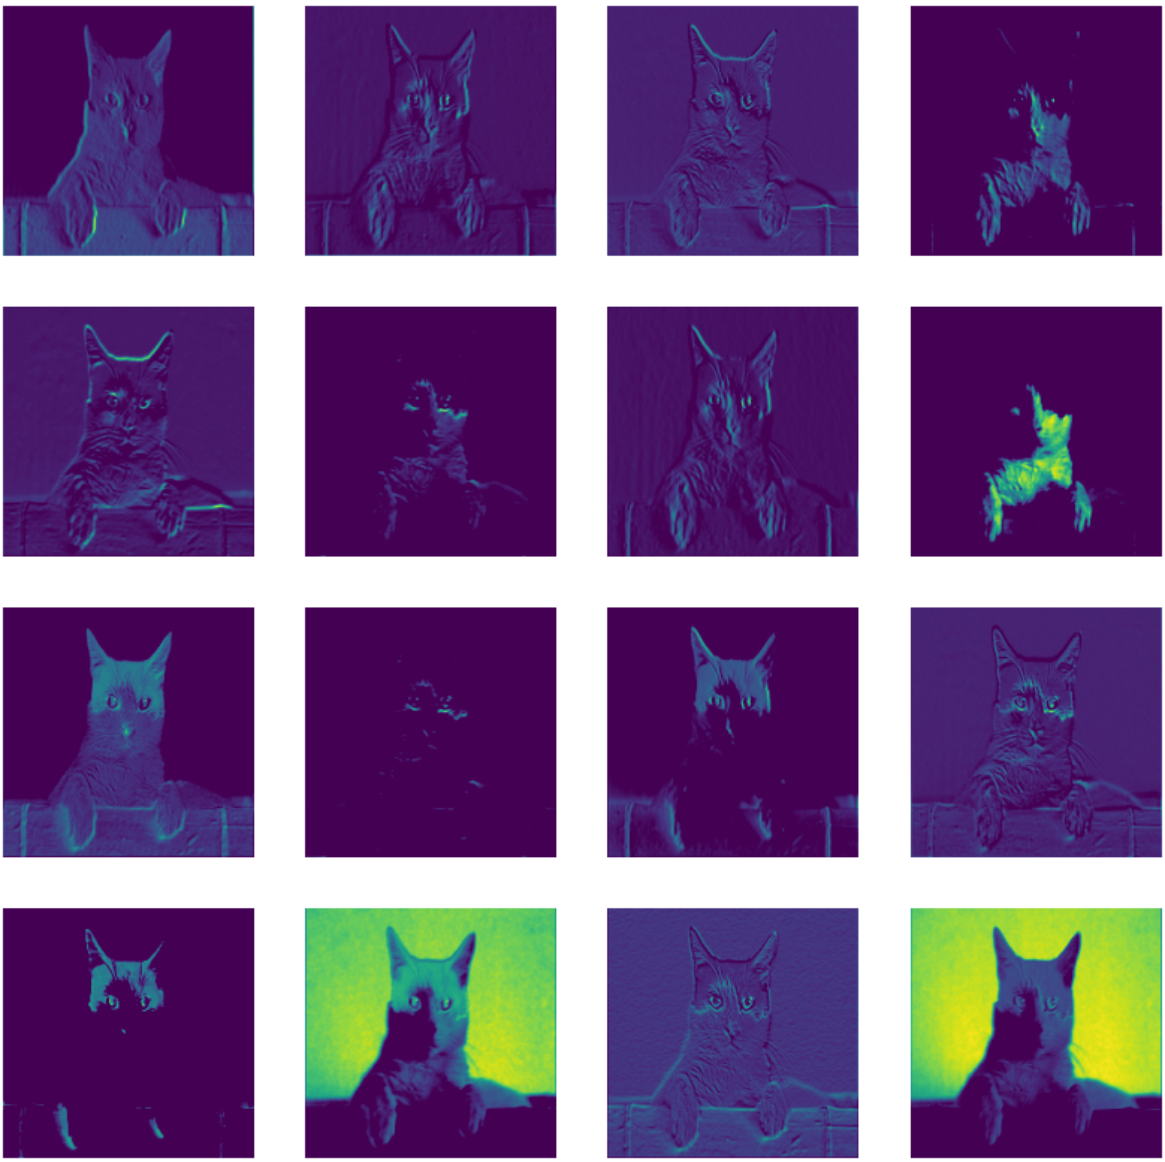

Figure 2: Visualize the output of first convolution.

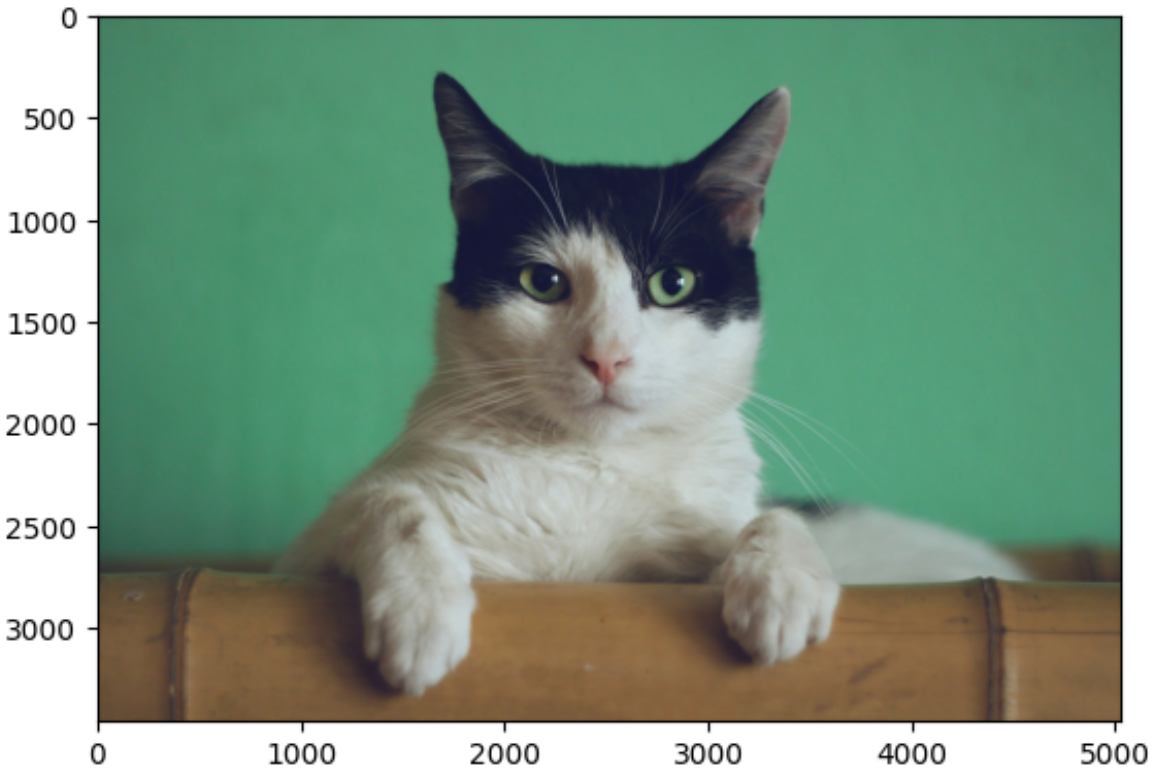

Figure 3: Origin input image of a cat.

★Question 5

Because VGG16 has too many parameters and 15 Scene is a relatively small dataset. If we directly train VGG16 on 15 Scene, it can cause significant overfitting.

★Question 6

Pre-training on ImageNet provides a strong baseline model that has already learned to recognize a wide range of features, such as edges, textures, shapes, and other complex patterns. This ability serves as a "universal" foundation for feature extraction, reducing training time and helping to prevent overfitting (as noted in question 5). This makes it a more effective approach for building a model that performs well on the 15 Scene dataset.

Question 7

First, the training dataset (ImageNet) may differ from the test dataset (15 Scene), so the features learned and the normalization parameters may not fully align. Second, since the task has changed, there might be a more suitable network architecture for extracting features specifically tailored for this classification task.

Question 8

In the early layers, the model primarily extracts basic graphical features like edges, textures, and light and shade. As more convolutional layers are added, the extracted features grow more complex, capturing shapes and partial structures of objects. In the final layers, the model can interpret complete information about the entire image, such as the object's class.

Question 9

Since the image is black and white, there is only one channel. We can make three copies of this channel after normalizing it and use it as the input to the RGB three channels

Question 10

A fully connected layer can certainly replace the SVM classifier. However, this approach has some drawbacks. For instance, it brings more parameters to be learned, and if the dataset is small, the large number of parameters can lead to overfitting.

Question 11

going further

2-b: Visualizing Neural Networks

★Question 1

From the following images, we can see that the primary object in the scene significantly influences the output of a trained image recognition neural network. In the second to fifth images, it is evident that the brightest part of the heat map corresponds to the location and contours of the animals. Similarly, in the first image depicting a haystack, the brightest regions correspond to the areas containing hay.

This demonstrates that the neural network effectively filters out background noise during computation, successfully identifying the main subject and producing accurate classifications.

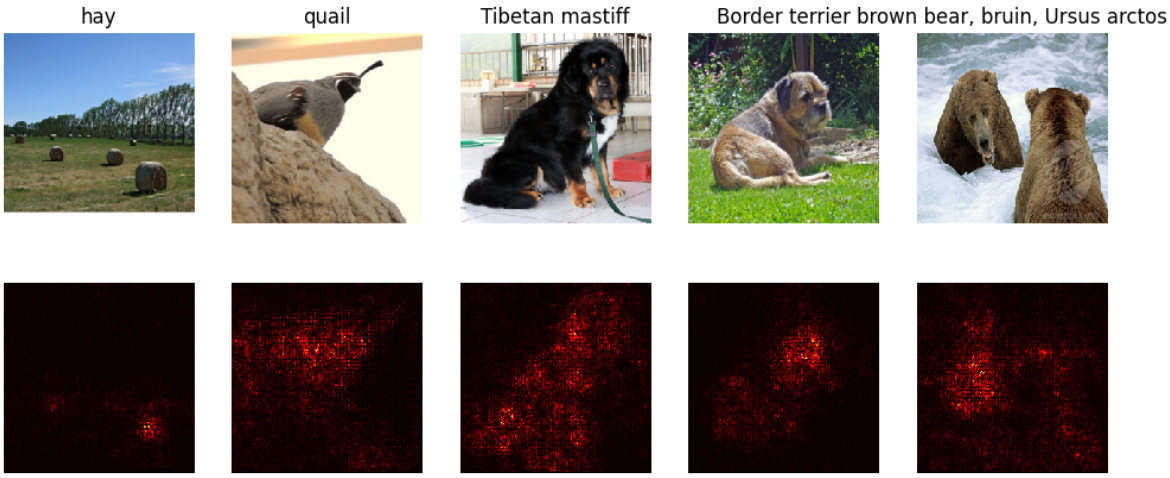

Figure 4: Saliency.

Question 2

1. **Gradient disappears:** In deep networks with highly non-linear activation functions, we may find the loss of gradient. This saturation may ignore some points that are also critical to the result.
2. **Class dependence:** Saliency maps only backward the gradient of a specific class. There may be other information that helps the network determine that the picture doesn't belong to another class, which is also important for the classification task, and is ignored.
3. **Lack of explanation:** Saliency maps only indicate which pixels contribute most to the model's prediction but do not provide the reason. It is purely mathematical and lacks explanation.

Question 3

It can be used for different purpose.

1. **Adversarial attack detection:** In the following part, we are already using it. This method can identify areas which is most sensitive to changes, enabling us only slightly modify the image to fool the classifier.
2. **Model debugging:**If the generated feature images are barely visible in relation to the object itself, this can indicate that our model for classification may not be effective.
3. **Class visualization:** As we will perform in Section 3, it can help us create the feature of a certain class.

Question 4

waiting....

★Question 5

The results are remarkable. Under conditions where the changes are nearly imperceptible to the naked eye, we successfully deceived the neural network into making an incorrect classification. This indicates that alterations to key pixels have a significant impact on the classification outcome. Furthermore, the original image label and the altered target label have a weak correlation, rather than being a simple confusion between similar categories, such as Labrador and Border Collie (both are dogs). This suggests that this approach is generalizable to such networks and has a low dependence on specific categories.

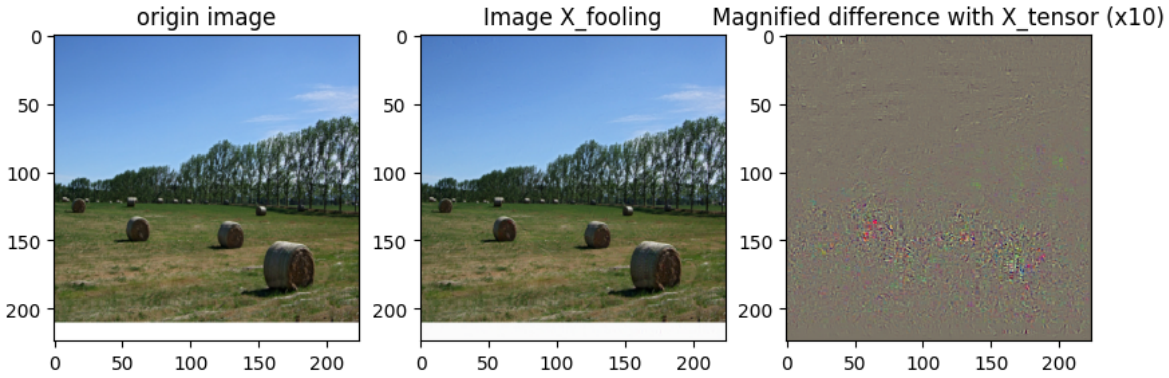

(a) From hay to stingray

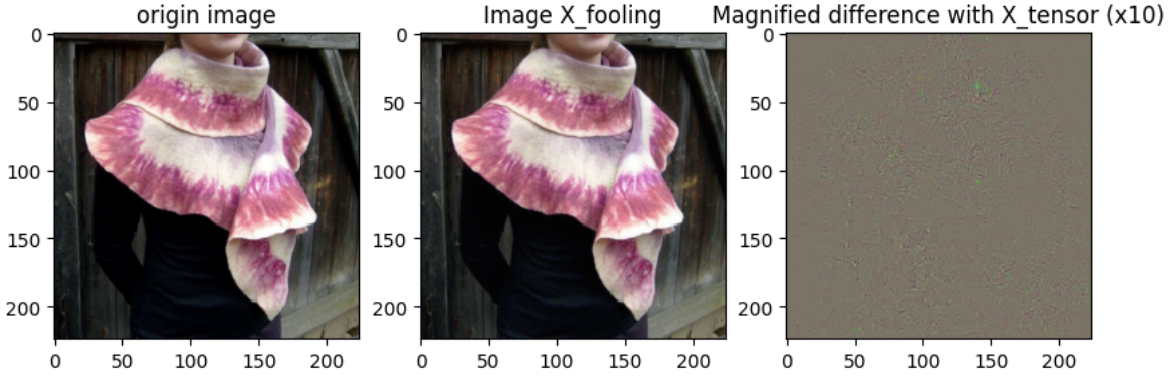

(b) From stole to stingray

Figure 5: Comparison of the original picture and the fooling image.

Question 6

1. **Copyright Protection:** This method can be employed as a means to protect copyright. There are artists who object to their creations being utilized as training material for AI systems, particularly without their consent. By employing this technique, they can modify their works prior to publication, thereby complicating the process for certain network to categorize them effectively and remain almost imperceptible to the human eye.

2. **Robustness Assurance and further study:**In certain domains, errors in recognition can lead to significant safety issues, such as in the field of autonomous driving. Therefore, it is imperative to investigate how to prevent situations where a small amount of noise causes the overall failure of image recognition. Such research also aids in a deeper understanding of the principles behind how neural networks recognize images.

3. **Class visualization:** As we will perform in Section 3, it can help us create the feature of a certain class.

Question 7

1. **Dependence on Full Model Access:** Dependence on Full Model Access: The method requires the computation of gradients, which necessitates access to the model's parameters. In practical scenarios, such as attacking black-box models, obtaining such access is often not feasible.
2. **Limited Generalization Ability:** The dependency on specific models means that adversarial examples generated may not perform effectively across other models.
3. **Computational Cost:** The requirement to compute gradients can lead to time consuming computation, for example, high-dimensional inputs or complex models.

Alternative ways:

- Fast Gradient Sign Method (FGSM), which is computationally efficient.
- Carlini & Wagner (C&W) Attack, which generates more imperceptible and effective examples.
- For black-box model, we can try to fine similar known model to generate fool examples. If we know noting about the model, we can try *Query-based Attacks* , e.g.:Zeroth-Order Optimization (ZOO),NES Attack...

★Question 8

The following figure 6 (figure 7) visualizes class tarantula (snail) from a irrelevant white noise (a picture of hay). In figure 6, without the disturb of original image, we can see more clearly that the network identifies this category by capturing features similar to spider legs or bodies. This visualization method indicates that what the network learns are not specific objects, but rather patterns of features, such as edges, textures, and shapes.

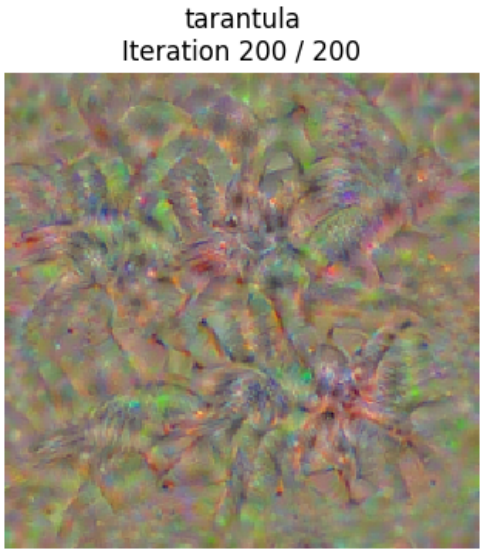

Figure 6: Noise to tarantula.

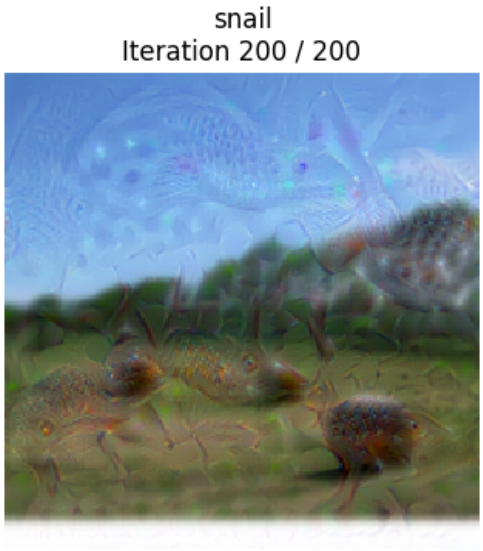

Figure 7: Hay to snail.

Question 9